GLEX-Alltoall: gather-scatter-based multi-leader All-to-all Communication on Multi-core Supercomputer

Jintao Peng[1], Jie Liu[2,*], Min Xie[3,*]

Changsha, China

Abstract

All-to-all communication is commonly used in parallel applications like FFT. In mordern supercomputers, there are multiple cores, NUMAs and network endpoints. These features bring much parallelism. However, there is no method which makes uses the parallelism to improve the all-to-all communication. In this paper, we introduce an optimized NUMA-aware multi-leader all-to-all library which explore the parallelism on network, CPU cores and overlap the intra- and inter-node communication. The results show that, compared to MPI, our library achieves up to 20x speedup. For application, our method achieves up to 1.75x speedup on peak performance for 16384 cores.

Keywords: Collective Communication, Multi-core processor, MPI all-to-all, RDMA, Shared Heap
2010 MSC: 00-01, 99-00

1. Introduction

Many parallel applications may suffer from global communication. Especially for communication-intensive applications, their time-to-solution and scalablily may be affected by global communication. Message Passing Interface (MPI) provides a set of commonly used collective communication. MPI_Alltoall is one of the collective communication where each process will send a different message to all processes. It is broadly used in some parallel applications like Fast Fourier Transform (FFT) [1] and some graph algorithms like MapReduce [2] and Breadth-first search (BFS) [3]. However, each time we double the processes, the all-to-all communication workload is quadrupled. On mordern supercomputers, network throughput has a linear relationship with the number of nodes. This brings great challenges to large-scale all-to-all communications.

For multiple-core processes, an effective way is node-aware all-to-all method [4]. It's replace a N nodes global all-to-all into N-1 times intra-node gather + local transpose + inter-node transpose + N-1 times intra-node gather. This method is very effective for small messages. Because, compared to original method, a node-aware all-to-all reduce the number of inter-node messages from $(M^N)^2$ to N^2 times (M is number of processers in each nodes). The size of the message is increased by M^2 times, which makes effective use of the network bandwidth. In the current supercomputer, a node has multiple CPU cores, NUMA and network endpoints. This architecture brings 4 kinds of parallism to optimize a node-aware all-to-all method:

(1) Multiple network endpoints can simultaneously process multiple communication requests.
(2) Processes in different NUMA can simultaneously access it local memory without contention.
(3) Multiple processes can simultaneously gather/scatter data and compose communication requests.
(4) Inter-node communication can be overlapped with intra-node communication.

As we known, no methods combine these parallism togather to improve a node-aware all-to-all collective communication.

In this paper, we proposed a multi-leader node-aware all-to-all method. It using multiple leaders on different NUMA which open the different network endpoints to gather/scatter data, compose communication requests, and transpose local matrix. It explore the parallelism existing in mordern multi-core processor with NUMA memory architecture and multi-port network. For intra-node gather/sactter, we proposed a shared-heap-based remote accessible memory which similar to intra-node MPI RMA. Inter-node communication is based on Remote Direct Memory Access (RDMA) which provide high throughput and low latency. The results show that, compared to MPI_alltoall, our implementation achieves up to 20x speedup and 4x speedup on average.

2. Related Work

From an algorithm perspective: Bruck algorithm [5] is commonly used for small message all-to-all. For mid size messages, isend-irecv algorithm is used. For large messages, linear shift exchange [6], pairwise exchange[7].

When considering the multi-core processors: Cache-oblivious MPI all-to-all (SH-NUMA-CO) based on morton order is proposed to minimize the cache miss rate [8]. For Infiniband and

*Corresponding author
[1]JintaoPengCS@gmail.com
[2]liujie@nudt.edu.cn
[3]xiemin@nudt.edu.cn

multi-core systems, a all-to-all collective (SA-orig) which based on shared memory aggregation techniques is proposed in [9]. For multi-rail QsNet SMP clusters, a shared memory and RDMA based all-to-all collectives (elan_alltoall) is proposed in [10]. For Intel Many Integrated Core (MIC) architecture, the re-routing scheme based all-to-all collective (PAIRWISE-SLR/BRUCK-SLR) is proposed in [11]. These works are direct related to our work. Table 1 shows the overall design-space for all-to-all collective on mulit-core processers.

When considering the network topology: A bandwidth-optimal all-to-all exchange is proposed for fat-tree network [12]. For torus network, a large scale all-to-all is proposed for Blue Gene/L Supercomputer [13]. A optimal schedule for all-to-all personalized communication is proposed for multiprocessor systems [14]. For Infiniband clusters, their is a topology aware all-to-all scheduler which lower the contention by adding extra communication steps [15].

When considering the network topology: A bandwidth-optimal all-to-all exchange is proposed for fat-tree network [12]. For torus network, a large scale all-to-all is proposed for Blue Gene/L Supercomputer [13]. A optimal schedule for all-to-all personalized communication is proposed for multiprocessor systems [14]. For Infiniband clusters, their is a topology aware all-to-all scheduler which lower the contention by adding extra communication steps [15].

When considering the network topology: A bandwidth-optimal all-to-all exchange is proposed for fat-tree network [12]. For torus network, a large scale all-to-all is proposed for Blue Gene/L Supercomputer [13]. A optimal schedule for all-to-all personalized communication is proposed for multiprocessor systems [14]. For Infiniband clusters, their is a topology aware all-to-all scheduler which lower the contention by adding extra communication steps [15].

When considering the network topology: A bandwidth-optimal all-to-all exchange is proposed for fat-tree network [12]. For torus network, a large scale all-to-all is proposed for Blue Gene/L Supercomputer [13]. A optimal schedule for all-to-all personalized communication is proposed for multiprocessor systems [14]. For Infiniband clusters, their is a topology aware all-to-all scheduler which lower the contention by adding extra communication steps [15].

When considering the network topology: A bandwidth-optimal all-to-all exchange is proposed for fat-tree network [12]. For torus network, a large scale all-to-all is proposed for Blue Gene/L Supercomputer [13]. A optimal schedule for all-to-all personalized communication is proposed for multiprocessor systems [14]. For Infiniband clusters, their is a topology aware all-to-all scheduler which lower the contention by adding extra communication steps [15].

When considering the network topology: A bandwidth-optimal all-to-all exchange is proposed for fat-tree network [12]. For torus network, a large scale all-to-all is proposed for Blue Gene/L Supercomputer [13]. A optimal schedule for all-to-all personalized communication is proposed for multiprocessor systems [14]. For Infiniband clusters, their is a topology aware all-to-all scheduler which lower the contention by adding extra communication steps [15].

When considering the network topology: A bandwidth-optimal all-to-all exchange is proposed for fat-tree network [12]. For torus network, a large scale all-to-all is proposed for Blue Gene/L Supercomputer [13]. A optimal schedule for all-to-all personalized communication is proposed for multiprocessor systems [14]. For Infiniband clusters, their is a topology aware all-to-all scheduler which lower the contention by adding extra communication steps [15].

When considering the network topology: A bandwidth-optimal all-to-all exchange is proposed for fat-tree network [12]. For torus network, a large scale all-to-all is proposed for Blue Gene/L Supercomputer [13]. A optimal schedule for all-to-all personalized communication is proposed for multiprocessor systems [14]. For Infiniband clusters, their is a topology aware all-to-all scheduler which lower the contention by adding extra communication steps [15].

When considering the network topology: A bandwidth-optimal all-to-all exchange is proposed for fat-tree network [12]. For torus network, a large scale all-to-all is proposed for Blue Gene/L Supercomputer [13]. A optimal schedule for all-to-all personalized communication is proposed for multiprocessor systems [14]. For Infiniband clusters, their is a topology aware all-to-all scheduler which lower the contention by adding extra communication steps [15].

When considering the network topology: A bandwidth-optimal all-to-all exchange is proposed for fat-tree network [12]. For torus network, a large scale all-to-all is proposed for Blue Gene/L Supercomputer [13]. A optimal schedule for all-to-all personalized communication is proposed for multiprocessor systems [14]. For Infiniband clusters, their is a topology aware all-to-all scheduler which lower the contention by adding extra communication steps [15].

3. Front matter

The author names and affiliations could be formatted in two ways:

(1) Group the authors per affiliation.

(2) Use footnotes to indicate the affiliations.

See the front matter of this document for examples. You are recommended to conform your choice to the journal you are submitting to.

4. Bibliography styles

There are various bibliography styles available. You can select the style of your choice in the preamble of this document. These styles are Elsevier styles based on standard styles like Harvard and Vancouver. Please use BibTEX to generate your bibliography and include DOIs whenever available.

Here are two sample references: [**? ?**].

References